

AI Governance, Risk, Compliance & Security

A structured reference covering 518 topics across 21 domains — from governance and model risk to generative AI security, LLM guardrails, and agentic AI oversight.

518

TOPICS

21

DOMAINS

64

PAGES

-
- | | |
|------------------------------------|---------------------------------------|
| 01 AI Governance | 12 LLM Guardrails |
| 02 AI Risk Management | 13 RAG Security and Governance |
| 03 Model Risk | 14 AI Agent and Agentic AI Governance |
| 04 Data Governance and Data Risk | 15 Privacy |
| 05 Responsible AI | 16 AI Testing and Assurance |
| 06 Bias and Fairness | 17 AI Audit |
| 07 Explainability and Transparency | 18 Third-Party and Vendor AI Risk |
| 08 AI Compliance | 19 AI Incident Management |
| 09 AI Regulations and Standards | 20 AI Lifecycle Management |
| 10 AI Security | 21 AI Security Tools and Frameworks |
| 11 Generative AI Security | |
-

Contents

518 topics organized into 21 domains, each explained in plain language with practical context.

#	DOMAIN	TOPICS	PAGE
01	AI Governance	48 topics	3
02	AI Risk Management	43 topics	8
03	Model Risk	27 topics	13
04	Data Governance and Data Risk	32 topics	16
05	Responsible AI	25 topics	20
06	Bias and Fairness	20 topics	23
07	Explainability and Transparency	14 topics	26
08	AI Compliance	17 topics	28
09	AI Regulations and Standards	27 topics	30
10	AI Security	29 topics	33
11	Generative AI Security	36 topics	36
12	LLM Guardrails	29 topics	40
13	RAG Security and Governance	20 topics	43
14	AI Agent and Agentic AI Governance	26 topics	45
15	Privacy	20 topics	48
16	AI Testing and Assurance	25 topics	51
17	AI Audit	17 topics	54
18	Third-Party and Vendor AI Risk	16 topics	56
19	AI Incident Management	15 topics	58
20	AI Lifecycle Management	15 topics	60
21	AI Security Tools and Frameworks	17 topics	62

01 AI Governance

S.NO.	TOPIC	EXPLANATION
1	AI Governance Fundamentals	AI Governance Fundamentals defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
2	AI Governance Frameworks	AI Governance Frameworks defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
3	AI Governance Operating Model	AI Governance Operating Model defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
4	AI Policies and Standards	AI Policies and Standards sets formal expectations for how AI should be designed, developed, acquired, deployed, and used. It translates organizational values, legal obligations, and risk appetite into practical requirements that teams can follow and audit.
5	AI Principles	AI Principles sets formal expectations for how AI should be designed, developed, acquired, deployed, and used. It translates organizational values, legal obligations, and risk appetite into practical requirements that teams can follow and audit.
6	Responsible AI	Responsible AI is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
7	Trustworthy AI	Trustworthy AI is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
8	AI Ethics	AI Ethics is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
9	AI Accountability	AI Accountability clarifies who is responsible for decisions, controls, approvals, and outcomes associated with AI systems or data. Clear accountability reduces gaps between technical, business, legal, risk, and security teams and supports effective escalation.
10	AI Transparency	AI Transparency helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.

S.NO.	TOPIC	EXPLANATION
11	AI Explainability	AI Explainability helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
12	AI Interpretability	AI Interpretability helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
13	AI Fairness	AI Fairness addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
14	AI Bias Management	AI Bias Management addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
15	Human Oversight	Human Oversight ensures that appropriate human judgment remains available for important AI decisions and actions. It defines when people must review, override, approve, stop, or escalate AI behavior, especially for high-impact or uncertain situations.
16	Human-in-the-Loop	Human-in-the-Loop ensures that appropriate human judgment remains available for important AI decisions and actions. It defines when people must review, override, approve, stop, or escalate AI behavior, especially for high-impact or uncertain situations.
17	AI Roles and Responsibilities	AI Roles and Responsibilities clarifies who is responsible for decisions, controls, approvals, and outcomes associated with AI systems or data. Clear accountability reduces gaps between technical, business, legal, risk, and security teams and supports effective escalation.
18	AI Governance Committees	AI Governance Committees defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
19	AI Center of Excellence	AI Center of Excellence defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
20	AI Lifecycle Governance	AI Lifecycle Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
21	AI Use Case Governance	AI Use Case Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.

S.NO.	TOPIC	EXPLANATION
22	AI Model Governance	AI Model Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
23	AI System Inventory	AI System Inventory provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
24	AI Asset Inventory	AI Asset Inventory provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
25	AI Use Case Inventory	AI Use Case Inventory provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
26	AI Risk Classification	AI Risk Classification focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
27	AI System Classification	AI System Classification is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
28	High-Risk AI Identification	High-Risk AI Identification focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
29	AI Impact Assessment	AI Impact Assessment is a structured assessment used to identify potential harms, affected stakeholders, and control requirements for an AI system. It records risks, safeguards, residual concerns, and approval decisions for governance and audit purposes.
30	Algorithmic Impact Assessment	Algorithmic Impact Assessment is a structured assessment used to identify potential harms, affected stakeholders, and control requirements for an AI system. It records risks, safeguards, residual concerns, and approval decisions for governance and audit purposes.
31	AI Approval Workflows	AI Approval Workflows is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

S.NO.	TOPIC	EXPLANATION
32	AI Change Management	AI Change Management is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
33	AI Documentation	AI Documentation provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
34	Model Cards	Model Cards provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
35	System Cards	System Cards provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
36	Data Sheets for Datasets	Data Sheets for Datasets provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
37	AI Decision Records	AI Decision Records provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
38	AI Audit Trails	AI Audit Trails provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
39	AI Monitoring and Reporting	AI Monitoring and Reporting provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
40	AI Governance Metrics and KPIs	AI Governance Metrics and KPIs defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
41	AI Vendor Governance	AI Vendor Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
42	Third-Party AI Governance	Third-Party AI Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
43	AI Procurement Governance	AI Procurement Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.

S.NO.	TOPIC	EXPLANATION
44	Shadow AI Governance	Shadow AI Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
45	Generative AI Governance	Generative AI Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
46	Agentic AI Governance	Agentic AI Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
47	Autonomous AI Governance	Autonomous AI Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
48	AI Decommissioning	AI Decommissioning is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.

02 AI Risk Management

S.NO.	TOPIC	EXPLANATION
1	AI Risk Management Fundamentals	AI Risk Management Fundamentals focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
2	AI Risk Management Frameworks	AI Risk Management Frameworks is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
3	AI Risk Identification	AI Risk Identification focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
4	AI Risk Assessment	AI Risk Assessment focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
5	AI Risk Analysis	AI Risk Analysis focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
6	AI Risk Evaluation	AI Risk Evaluation focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
7	AI Risk Prioritization	AI Risk Prioritization focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
8	AI Risk Treatment	AI Risk Treatment focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
9	AI Risk Mitigation	AI Risk Mitigation focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.

S.NO.	TOPIC	EXPLANATION
10	AI Risk Acceptance	AI Risk Acceptance focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
11	AI Risk Registers	AI Risk Registers focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
12	AI Control Mapping	AI Control Mapping is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
13	Inherent Risk	Inherent Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
14	Residual Risk	Residual Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
15	Operational Risk	Operational Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
16	Model Risk	Model Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
17	Data Risk	Data Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
18	Privacy Risk	Privacy Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
19	Security Risk	Security Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.

S.NO.	TOPIC	EXPLANATION
20	Ethical Risk	Ethical Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
21	Legal Risk	Legal Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
22	Regulatory Risk	Regulatory Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
23	Reputational Risk	Reputational Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
24	Financial Risk	Financial Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
25	Third-Party Risk	Third-Party Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
26	Supply Chain Risk	Supply Chain Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
27	Concentration Risk	Concentration Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
28	Business Continuity Risk	Business Continuity Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
29	AI System Failure Risk	AI System Failure Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.

S.NO.	TOPIC	EXPLANATION
30	Automation Risk	Automation Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
31	Human Dependency Risk	Human Dependency Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
32	Over-Reliance on AI	Over-Reliance on AI is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
33	AI Misuse Risk	AI Misuse Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
34	AI Abuse Risk	AI Abuse Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
35	AI Safety Risk	AI Safety Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
36	Systemic AI Risk	Systemic AI Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
37	Emerging AI Risk	Emerging AI Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
38	Generative AI Risk	Generative AI Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
39	Agentic AI Risk	Agentic AI Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.

S.NO. TOPIC	EXPLANATION
40 AI Risk Monitoring	AI Risk Monitoring focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
41 Key Risk Indicators	Key Risk Indicators focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
42 AI Incident Management	AI Incident Management is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
43 AI Risk Reporting	AI Risk Reporting focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.

03 Model Risk

S.NO.	TOPIC	EXPLANATION
1	Model Risk Management	Model Risk Management focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
2	Model Development Risk	Model Development Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
3	Model Validation	Model Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
4	Independent Model Validation	Independent Model Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
5	Model Testing	Model Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
6	Model Verification	Model Verification evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
7	Model Performance Risk	Model Performance Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
8	Model Accuracy	Model Accuracy concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
9	Model Robustness	Model Robustness concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.

S.NO.	TOPIC	EXPLANATION
10	Model Reliability	Model Reliability concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
11	Model Stability	Model Stability concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
12	Model Drift	Model Drift concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
13	Data Drift	Data Drift concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
14	Concept Drift	Concept Drift concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
15	Model Decay	Model Decay concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
16	Model Calibration	Model Calibration concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
17	Model Uncertainty	Model Uncertainty concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
18	Model Hallucination	Model Hallucination concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
19	Model Overfitting	Model Overfitting concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.

S.NO.	TOPIC	EXPLANATION
20	Model Underfitting	Model Underfitting concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
21	Model Generalization	Model Generalization concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
22	Model Benchmarking	Model Benchmarking evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
23	Model Versioning	Model Versioning provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
24	Model Lineage	Model Lineage provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
25	Model Monitoring	Model Monitoring provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
26	Model Retraining Governance	Model Retraining Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
27	Model Retirement	Model Retirement is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.

04 Data Governance and Data Risk

S.NO.	TOPIC	EXPLANATION
1	Data Governance	Data Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
2	Data Quality	Data Quality is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
3	Data Accuracy	Data Accuracy concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
4	Data Completeness	Data Completeness is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
5	Data Consistency	Data Consistency is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
6	Data Integrity	Data Integrity is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
7	Data Provenance	Data Provenance is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
8	Data Lineage	Data Lineage provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
9	Data Ownership	Data Ownership clarifies who is responsible for decisions, controls, approvals, and outcomes associated with AI systems or data. Clear accountability reduces gaps between technical, business, legal, risk, and security teams and supports effective escalation.
10	Data Stewardship	Data Stewardship clarifies who is responsible for decisions, controls, approvals, and outcomes associated with AI systems or data. Clear accountability reduces gaps between technical, business, legal, risk, and security teams and supports effective escalation.

S.NO.	TOPIC	EXPLANATION
11	Data Classification	Data Classification is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
12	Data Minimization	Data Minimization is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
13	Data Retention	Data Retention is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
14	Data Residency	Data Residency is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
15	Data Sovereignty	Data Sovereignty is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
16	Data Access Control	Data Access Control protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
17	Sensitive Data Management	Sensitive Data Management is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
18	Personally Identifiable Information	Personally Identifiable Information is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
19	Sensitive Personal Data	Sensitive Personal Data protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
20	Training Data Governance	Training Data Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.

S.NO.	TOPIC	EXPLANATION
21	Validation Data Governance	Validation Data Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
22	Test Data Governance	Test Data Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
23	Synthetic Data Governance	Synthetic Data Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
24	Data Labeling Governance	Data Labeling Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
25	Data Bias	Data Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
26	Dataset Representativeness	Dataset Representativeness is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
27	Copyrighted Training Data	Copyrighted Training Data is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
28	Data Licensing	Data Licensing is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
29	Data Consent	Data Consent is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
30	Data Poisoning	Data Poisoning is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.

S.NO. TOPIC	EXPLANATION
31 Data Leakage	Data Leakage is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
32 Data Exfiltration	Data Exfiltration is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

05 Responsible AI

S.NO.	TOPIC	EXPLANATION
1	Responsible AI Principles	Responsible AI Principles sets formal expectations for how AI should be designed, developed, acquired, deployed, and used. It translates organizational values, legal obligations, and risk appetite into practical requirements that teams can follow and audit.
2	Fairness	Fairness addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
3	Bias	Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
4	Non-Discrimination	Non-Discrimination addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
5	Transparency	Transparency helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
6	Explainability	Explainability helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
7	Accountability	Accountability clarifies who is responsible for decisions, controls, approvals, and outcomes associated with AI systems or data. Clear accountability reduces gaps between technical, business, legal, risk, and security teams and supports effective escalation.
8	Privacy	Privacy protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
9	Security	Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
10	Safety	Safety is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

S.NO.	TOPIC	EXPLANATION
11	Reliability	Reliability concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
12	Robustness	Robustness concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
13	Inclusiveness	Inclusiveness is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
14	Accessibility	Accessibility is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
15	Human-Centric AI	Human-Centric AI is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
16	Human Agency	Human Agency ensures that appropriate human judgment remains available for important AI decisions and actions. It defines when people must review, override, approve, stop, or escalate AI behavior, especially for high-impact or uncertain situations.
17	Human Oversight	Human Oversight ensures that appropriate human judgment remains available for important AI decisions and actions. It defines when people must review, override, approve, stop, or escalate AI behavior, especially for high-impact or uncertain situations.
18	Contestability	Contestability concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
19	Traceability	Traceability is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
20	Sustainability	Sustainability is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

S.NO. TOPIC	EXPLANATION
21 Social Impact	Social Impact is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
22 Environmental Impact	Environmental Impact is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
23 Ethical AI Design	Ethical AI Design is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
24 Responsible AI Assessments	Responsible AI Assessments is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
25 Responsible AI Controls	Responsible AI Controls is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

06 Bias and Fairness

S.NO.	TOPIC	EXPLANATION
1	Algorithmic Bias	Algorithmic Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
2	Data Bias	Data Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
3	Historical Bias	Historical Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
4	Sampling Bias	Sampling Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
5	Selection Bias	Selection Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
6	Measurement Bias	Measurement Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
7	Representation Bias	Representation Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
8	Label Bias	Label Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
9	Confirmation Bias	Confirmation Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.

S.NO.	TOPIC	EXPLANATION
10	Automation Bias	Automation Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
11	Proxy Bias	Proxy Bias addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
12	Fairness Metrics	Fairness Metrics addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
13	Demographic Parity	Demographic Parity addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
14	Equal Opportunity	Equal Opportunity addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
15	Equalized Odds	Equalized Odds addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
16	Disparate Impact	Disparate Impact addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
17	Bias Detection	Bias Detection addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
18	Bias Testing	Bias Testing addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
19	Bias Mitigation	Bias Mitigation addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.

S.NO.	TOPIC	EXPLANATION
20	Fairness Monitoring	Fairness Monitoring addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.

07 Explainability and Transparency

S.NO.	TOPIC	EXPLANATION
1	Explainable AI	Explainable AI helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
2	Interpretable AI	Interpretable AI helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
3	Global Explainability	Global Explainability helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
4	Local Explainability	Local Explainability helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
5	Feature Importance	Feature Importance helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
6	SHAP	SHAP helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
7	LIME	LIME helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
8	Counterfactual Explanations	Counterfactual Explanations helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
9	Decision Transparency	Decision Transparency helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.

S.NO.	TOPIC	EXPLANATION
10	Model Transparency	Model Transparency helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
11	Data Transparency	Data Transparency helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
12	User Disclosure	User Disclosure is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
13	AI-Generated Content Disclosure	AI-Generated Content Disclosure is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
14	Explainability Documentation	Explainability Documentation helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.

08 AI Compliance

S.NO.	TOPIC	EXPLANATION
1	AI Compliance Management	AI Compliance Management ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
2	Regulatory Compliance	Regulatory Compliance ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
3	Legal Compliance	Legal Compliance ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
4	Internal Policy Compliance	Internal Policy Compliance sets formal expectations for how AI should be designed, developed, acquired, deployed, and used. It translates organizational values, legal obligations, and risk appetite into practical requirements that teams can follow and audit.
5	Compliance Controls	Compliance Controls ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
6	Compliance Monitoring	Compliance Monitoring provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
7	Compliance Testing	Compliance Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
8	Compliance Evidence	Compliance Evidence ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
9	Compliance Reporting	Compliance Reporting provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
10	Regulatory Mapping	Regulatory Mapping ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.

S.NO.	TOPIC	EXPLANATION
11	Regulatory Gap Assessment	Regulatory Gap Assessment ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
12	AI Compliance Audits	AI Compliance Audits provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
13	Control Testing	Control Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
14	Continuous Compliance	Continuous Compliance ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
15	Compliance Automation	Compliance Automation ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
16	Record Keeping	Record Keeping provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
17	Documentation Requirements	Documentation Requirements provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.

09 AI Regulations and Standards

S.NO.	TOPIC	EXPLANATION
1	EU AI Act	A European Union regulation that establishes risk-based requirements for AI systems, including obligations for prohibited, high-risk, transparency, and general-purpose AI use cases. It helps organizations determine what controls, documentation, monitoring, and accountability measures are legally required.
2	GDPR	The General Data Protection Regulation governs how personal data is collected, processed, stored, and protected in the EU and EEA. For AI systems, it is especially relevant to lawful processing, transparency, profiling, automated decision-making, and data-subject rights.
3	Digital Services Act	Digital Services Act is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
4	Digital Markets Act	Digital Markets Act is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
5	NIST AI Risk Management Framework	The NIST AI RMF provides a voluntary framework for identifying, measuring, managing, and governing risks from AI systems. It organizes activities around Govern, Map, Measure, and Manage functions across the AI lifecycle.
6	NIST Cybersecurity Framework	The NIST Cybersecurity Framework provides a structured approach for managing cybersecurity risk using functions such as Govern, Identify, Protect, Detect, Respond, and Recover. It can be applied to protect AI models, data, infrastructure, APIs, and supporting services.
7	ISO/IEC 42001	ISO/IEC 42001 is a management-system standard for establishing, implementing, maintaining, and continually improving an AI Management System. It helps organizations formalize governance, risk management, responsibilities, controls, and oversight for AI.
8	ISO/IEC 23894	ISO/IEC 23894 provides guidance for managing risks associated with artificial intelligence. It supports systematic identification, assessment, treatment, monitoring, and communication of AI risks throughout the lifecycle.
9	ISO/IEC 27001	ISO/IEC 27001 defines requirements for an Information Security Management System. In AI environments, it supports governance of access control, asset management, incident handling, supplier risk, and protection of models and data.
10	ISO/IEC 27701	ISO/IEC 27701 extends information-security management with privacy information management controls. It is useful for governing personal data used in AI training, inference, analytics, and automated decision processes.

S.NO.	TOPIC	EXPLANATION
11	ISO/IEC 27017	ISO/IEC 27017 provides security-control guidance for cloud services and cloud customers. It is relevant when AI workloads, models, vector databases, or data pipelines are hosted in cloud environments.
12	ISO/IEC 27018	ISO/IEC 27018 provides privacy controls for protecting personally identifiable information in public-cloud services. It helps organizations manage privacy obligations when AI services process personal information in cloud platforms.
13	ISO/IEC 22989	ISO/IEC 22989 defines foundational AI concepts and terminology. It helps teams use consistent language when discussing AI systems, stakeholders, capabilities, risks, and lifecycle activities.
14	ISO/IEC 23053	ISO/IEC 23053 describes a framework for AI systems that use machine learning. It provides a common architectural view of ML components, data flows, model development, and operational interactions.
15	OECD AI Principles	The OECD AI Principles promote innovative and trustworthy AI that respects human rights and democratic values. They emphasize inclusive growth, human-centered values, transparency, robustness, security, safety, and accountability.
16	UNESCO Recommendation on AI Ethics	UNESCO's Recommendation on the Ethics of Artificial Intelligence provides global guidance for developing and using AI in ways that protect human rights and social well-being. It addresses governance, fairness, transparency, privacy, sustainability, and human oversight.
17	IEEE AI Standards	IEEE AI Standards sets formal expectations for how AI should be designed, developed, acquired, deployed, and used. It translates organizational values, legal obligations, and risk appetite into practical requirements that teams can follow and audit.
18	COBIT	COBIT is an enterprise governance framework for managing information and technology. It can be extended to AI governance by defining objectives, controls, accountability, risk management, assurance, and performance measures.
19	COSO	COSO provides frameworks for enterprise risk management and internal control. Organizations can use COSO concepts to integrate AI risks into governance, risk appetite, control design, monitoring, and assurance.
20	SOC 2	SOC 2 evaluates controls related to security, availability, processing integrity, confidentiality, and privacy for service organizations. AI service providers may use SOC 2 controls to demonstrate operational and security assurance to customers.
21	PCI DSS	PCI DSS defines security requirements for environments that store, process, or transmit payment-card data. AI applications handling payment information must ensure that data flows, logs, prompts, and integrations do not expose cardholder data.
22	HIPAA	HIPAA establishes privacy and security requirements for protected health information in the United States. AI used in healthcare must protect patient data and apply appropriate safeguards, access controls, and disclosure limitations.

S.NO.	TOPIC	EXPLANATION
23	CCPA/CPRA	CCPA and CPRA provide California consumers with privacy rights over personal information and impose obligations on covered businesses. AI systems may need controls for data access, deletion, disclosure, profiling, and sensitive-information use.
24	DORA	DORA is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
25	NIS2	NIS2 is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
26	SEC AI-related Requirements	SEC AI-related Requirements is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
27	Sector-Specific AI Regulations	Sector-Specific AI Regulations is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

10 AI Security

S.NO.	TOPIC	EXPLANATION
1	AI Security Fundamentals	AI Security Fundamentals protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
2	AI Security Governance	AI Security Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
3	Secure AI Development Lifecycle	Secure AI Development Lifecycle is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
4	AI Threat Modeling	AI Threat Modeling protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
5	AI Attack Surface	AI Attack Surface protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
6	AI Security Architecture	AI Security Architecture protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
7	AI Security Controls	AI Security Controls protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
8	AI Security Testing	AI Security Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
9	AI Red Teaming	AI Red Teaming evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.

S.NO.	TOPIC	EXPLANATION
10	AI Penetration Testing	AI Penetration Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
11	AI Vulnerability Management	AI Vulnerability Management protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
12	AI Incident Response	AI Incident Response is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
13	AI Security Monitoring	AI Security Monitoring provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
14	AI Security Logging	AI Security Logging protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
15	AI Security Auditing	AI Security Auditing protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
16	Secure Model Deployment	Secure Model Deployment is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
17	Secure Model Hosting	Secure Model Hosting is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
18	Model Access Control	Model Access Control protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
19	Model Authentication	Model Authentication protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.

S.NO.	TOPIC	EXPLANATION
20	Model Authorization	Model Authorization protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
21	Secrets Management	Secrets Management protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
22	Key Management	Key Management protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
23	Encryption	Encryption protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
24	Network Security	Network Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
25	API Security	API Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
26	Endpoint Security	Endpoint Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
27	Cloud AI Security	Cloud AI Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
28	Container Security	Container Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
29	AI Supply Chain Security	AI Supply Chain Security manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.

11 Generative AI Security

S.NO.	TOPIC	EXPLANATION
1	Generative AI Security	Generative AI Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
2	LLM Security	LLM Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
3	OWASP Top 10 for LLM Applications	OWASP Top 10 for LLM Applications is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
4	Prompt Injection	Prompt Injection is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
5	Direct Prompt Injection	Direct Prompt Injection is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
6	Indirect Prompt Injection	Indirect Prompt Injection is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
7	Jailbreaking	Jailbreaking is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
8	Prompt Leaking	Prompt Leaking is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
9	System Prompt Extraction	System Prompt Extraction is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.

S.NO.	TOPIC	EXPLANATION
10	Sensitive Information Disclosure	Sensitive Information Disclosure is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
11	Insecure Output Handling	Insecure Output Handling is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
12	Excessive Agency	Excessive Agency is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
13	Overprivileged AI Agents	Overprivileged AI Agents controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.
14	Insecure Plugin Design	Insecure Plugin Design is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
15	Insecure Tool Use	Insecure Tool Use is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
16	Model Denial of Service	Model Denial of Service is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
17	Model Theft	Model Theft is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
18	Model Extraction	Model Extraction is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
19	Training Data Extraction	Training Data Extraction is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.

S.NO.	TOPIC	EXPLANATION
20	Membership Inference	Membership Inference is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
21	Model Inversion	Model Inversion is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
22	Adversarial Inputs	Adversarial Inputs is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
23	Adversarial Examples	Adversarial Examples is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
24	Data Poisoning	Data Poisoning is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
25	Model Poisoning	Model Poisoning is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
26	Backdoor Attacks	Backdoor Attacks is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
27	Supply Chain Attacks	Supply Chain Attacks manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
28	Retrieval Poisoning	Retrieval Poisoning is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
29	RAG Poisoning	RAG Poisoning is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.

S.NO.	TOPIC	EXPLANATION
30	Vector Database Security	Vector Database Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
31	Embedding Security	Embedding Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
32	Context Window Attacks	Context Window Attacks is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
33	Context Manipulation	Context Manipulation is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
34	Token Smuggling	Token Smuggling is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
35	Tool Injection	Tool Injection is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
36	Agent Hijacking	Agent Hijacking is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.

12 LLM Guardrails

S.NO.	TOPIC	EXPLANATION
1	Input Guardrails	Input Guardrails is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
2	Output Guardrails	Output Guardrails is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
3	Prompt Guardrails	Prompt Guardrails is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
4	Content Filtering	Content Filtering is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
5	Toxicity Detection	Toxicity Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
6	Profanity Detection	Profanity Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
7	Hate Speech Detection	Hate Speech Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
8	Violence Detection	Violence Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
9	Self-Harm Detection	Self-Harm Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.

S.NO.	TOPIC	EXPLANATION
10	Sexual Content Filtering	Sexual Content Filtering is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
11	PII Detection	PII Detection protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
12	PII Masking	PII Masking protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
13	PII Redaction	PII Redaction protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
14	Secret Detection	Secret Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
15	Prompt Injection Detection	Prompt Injection Detection is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
16	Jailbreak Detection	Jailbreak Detection is an adversarial technique that attempts to manipulate an LLM or agent into ignoring intended instructions, exposing protected information, or performing unauthorized actions. Defenses include instruction isolation, filtering, tool permission controls, validation, monitoring, and adversarial testing.
17	Hallucination Detection	Hallucination Detection concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
18	Groundedness Checking	Groundedness Checking is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
19	Factuality Checking	Factuality Checking is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.

S.NO.	TOPIC	EXPLANATION
20	Relevance Checking	Relevance Checking is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
21	Topic Restriction	Topic Restriction is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
22	Policy Enforcement	Policy Enforcement sets formal expectations for how AI should be designed, developed, acquired, deployed, and used. It translates organizational values, legal obligations, and risk appetite into practical requirements that teams can follow and audit.
23	Schema Validation	Schema Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
24	Output Validation	Output Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
25	Structured Output Enforcement	Structured Output Enforcement is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
26	Allow Lists	Allow Lists is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
27	Block Lists	Block Lists is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
28	Rate Limiting	Rate Limiting is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
29	Human Escalation	Human Escalation is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.

13 RAG Security and Governance

S.NO.	TOPIC	EXPLANATION
1	RAG Governance	RAG Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
2	RAG Security	RAG Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
3	Document Access Control	Document Access Control protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
4	Knowledge Base Governance	Knowledge Base Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
5	Retrieval Authorization	Retrieval Authorization protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
6	Document-Level Permissions	Document-Level Permissions addresses governance, quality, or security within retrieval-augmented generation systems. Controls should ensure that retrieved content is authorized, trustworthy, current, traceable, relevant, and protected from poisoning or unauthorized disclosure.
7	Chunk-Level Permissions	Chunk-Level Permissions addresses governance, quality, or security within retrieval-augmented generation systems. Controls should ensure that retrieved content is authorized, trustworthy, current, traceable, relevant, and protected from poisoning or unauthorized disclosure.
8	Vector Database Governance	Vector Database Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
9	Vector Database Access Control	Vector Database Access Control protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
10	Embedding Privacy	Embedding Privacy protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.

S.NO.	TOPIC	EXPLANATION
11	Retrieval Poisoning	Retrieval Poisoning is an AI-specific attack or abuse scenario targeting the confidentiality, integrity, or behavior of models and their data. Mitigation typically combines access controls, rate limits, provenance checks, validation, anomaly detection, secure pipelines, and red-team testing.
12	Malicious Documents	Malicious Documents addresses governance, quality, or security within retrieval-augmented generation systems. Controls should ensure that retrieved content is authorized, trustworthy, current, traceable, relevant, and protected from poisoning or unauthorized disclosure.
13	Source Validation	Source Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
14	Source Attribution	Source Attribution addresses governance, quality, or security within retrieval-augmented generation systems. Controls should ensure that retrieved content is authorized, trustworthy, current, traceable, relevant, and protected from poisoning or unauthorized disclosure.
15	Citation Validation	Citation Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
16	Groundedness	Groundedness is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
17	Retrieval Relevance	Retrieval Relevance addresses governance, quality, or security within retrieval-augmented generation systems. Controls should ensure that retrieved content is authorized, trustworthy, current, traceable, relevant, and protected from poisoning or unauthorized disclosure.
18	Knowledge Freshness	Knowledge Freshness addresses governance, quality, or security within retrieval-augmented generation systems. Controls should ensure that retrieved content is authorized, trustworthy, current, traceable, relevant, and protected from poisoning or unauthorized disclosure.
19	Knowledge Versioning	Knowledge Versioning provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
20	Document Lineage	Document Lineage provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.

14 AI Agent and Agentic AI Governance

S.NO.	TOPIC	EXPLANATION
1	AI Agent Governance	AI Agent Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
2	Agentic AI Governance	Agentic AI Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
3	Autonomous Agent Risk	Autonomous Agent Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
4	Agent Identity	Agent Identity controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.
5	Agent Authentication	Agent Authentication protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
6	Agent Authorization	Agent Authorization protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
7	Agent Permissions	Agent Permissions controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.
8	Tool Governance	Tool Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
9	Tool Access Control	Tool Access Control protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
10	Least Privilege	Least Privilege protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.

S.NO.	TOPIC	EXPLANATION
11	Excessive Agency	Excessive Agency is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
12	Agent Memory Governance	Agent Memory Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
13	Agent Memory Security	Agent Memory Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
14	Multi-Agent Governance	Multi-Agent Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.
15	Agent-to-Agent Security	Agent-to-Agent Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
16	Agent Communication Security	Agent Communication Security protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
17	Agent Planning Risk	Agent Planning Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
18	Agent Execution Risk	Agent Execution Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
19	Human Approval Gates	Human Approval Gates ensures that appropriate human judgment remains available for important AI decisions and actions. It defines when people must review, override, approve, stop, or escalate AI behavior, especially for high-impact or uncertain situations.
20	Transaction Limits	Transaction Limits controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.

S.NO.	TOPIC	EXPLANATION
21	Agent Sandboxing	Agent Sandboxing protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
22	Agent Monitoring	Agent Monitoring provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
23	Agent Audit Trails	Agent Audit Trails provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
24	Kill Switches	Kill Switches controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.
25	Agent Rollback	Agent Rollback controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.
26	Agent Incident Response	Agent Incident Response controls the behavior and permissions of AI agents that can plan, remember context, use tools, communicate, or take actions. Effective governance limits autonomy appropriately, enforces identity and least privilege, records actions, and provides human intervention or shutdown mechanisms.

15 Privacy

S.NO.	TOPIC	EXPLANATION
1	AI Privacy	AI Privacy protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
2	Privacy by Design	Privacy by Design protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
3	Privacy by Default	Privacy by Default protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
4	Consent Management	Consent Management protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
5	Purpose Limitation	Purpose Limitation is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
6	Data Minimization	Data Minimization is a data-governance control that ensures information used by AI is appropriate, trustworthy, lawful, and manageable throughout its lifecycle. It should be supported by defined ownership, documented rules, technical checks, access controls, and evidence of compliance.
7	Data Anonymization	Data Anonymization protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
8	Data Pseudonymization	Data Pseudonymization protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
9	Differential Privacy	Differential Privacy protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.

S.NO.	TOPIC	EXPLANATION
10	Privacy-Preserving AI	Privacy-Preserving AI protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
11	Federated Learning	Federated Learning protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
12	Secure Multi-Party Computation	Secure Multi-Party Computation protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
13	Homomorphic Encryption	Homomorphic Encryption protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
14	Right to Access	Right to Access is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
15	Right to Erasure	Right to Erasure is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
16	Right to Explanation	Right to Explanation is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
17	Data Subject Rights	Data Subject Rights is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
18	Cross-Border Data Transfers	Cross-Border Data Transfers protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
19	Privacy Impact Assessment	Privacy Impact Assessment is a structured assessment used to identify potential harms, affected stakeholders, and control requirements for an AI system. It records risks, safeguards, residual concerns, and approval decisions for governance and audit purposes.

S.NO.	TOPIC	EXPLANATION
20	Data Protection Impact Assessment	Data Protection Impact Assessment is a structured assessment used to identify potential harms, affected stakeholders, and control requirements for an AI system. It records risks, safeguards, residual concerns, and approval decisions for governance and audit purposes.

16 AI Testing and Assurance

S.NO.	TOPIC	EXPLANATION
1	AI Assurance	AI Assurance evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
2	AI Validation	AI Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
3	AI Verification	AI Verification evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
4	Functional Testing	Functional Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
5	Performance Testing	Performance Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
6	Robustness Testing	Robustness Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
7	Fairness Testing	Fairness Testing addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
8	Bias Testing	Bias Testing addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
9	Security Testing	Security Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.

S.NO.	TOPIC	EXPLANATION
10	Privacy Testing	Privacy Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
11	Safety Testing	Safety Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
12	Explainability Testing	Explainability Testing helps stakeholders understand how an AI system reaches outputs, what information influences decisions, and what limitations apply. It supports trust, validation, user communication, regulatory review, and investigation when an AI decision is challenged.
13	Hallucination Testing	Hallucination Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
14	Toxicity Testing	Toxicity Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
15	Prompt Injection Testing	Prompt Injection Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
16	Jailbreak Testing	Jailbreak Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
17	Red Teaming	Red Teaming evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
18	Adversarial Testing	Adversarial Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
19	Stress Testing	Stress Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.

S.NO.	TOPIC	EXPLANATION
20	Regression Testing	Regression Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
21	Benchmark Testing	Benchmark Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
22	Scenario Testing	Scenario Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
23	Continuous Evaluation	Continuous Evaluation is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
24	LLM Evaluation	LLM Evaluation is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
25	AI Evaluation Metrics	AI Evaluation Metrics provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.

17 AI Audit

S.NO.	TOPIC	EXPLANATION
1	AI Auditing	AI Auditing provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
2	Internal AI Audit	Internal AI Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
3	External AI Audit	External AI Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
4	Algorithm Audit	Algorithm Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
5	Model Audit	Model Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
6	Data Audit	Data Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
7	Security Audit	Security Audit protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
8	Privacy Audit	Privacy Audit protects individuals and sensitive information when data is collected, transformed, transferred, or processed by AI systems. Controls may include minimization, masking, access restrictions, encryption, legal review, privacy-enhancing technologies, and auditable handling procedures.
9	Bias Audit	Bias Audit addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
10	Compliance Audit	Compliance Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
11	Responsible AI Audit	Responsible AI Audit provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.

S.NO.	TOPIC	EXPLANATION
12	Audit Evidence	Audit Evidence provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
13	Audit Trails	Audit Trails provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
14	Audit Logging	Audit Logging provides structured assurance that AI controls, processes, and evidence meet defined requirements. It examines design and operating effectiveness, records findings, and tracks remediation to improve governance and compliance.
15	Control Effectiveness Testing	Control Effectiveness Testing evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
16	Findings Management	Findings Management is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
17	Remediation Tracking	Remediation Tracking is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.

18 Third-Party and Vendor AI Risk

S.NO.	TOPIC	EXPLANATION
1	Third-Party AI Risk	Third-Party AI Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
2	Vendor Risk Management	Vendor Risk Management focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
3	AI Vendor Due Diligence	AI Vendor Due Diligence manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
4	Vendor Security Assessment	Vendor Security Assessment manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
5	Vendor Privacy Assessment	Vendor Privacy Assessment manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
6	Vendor Compliance Assessment	Vendor Compliance Assessment manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
7	Foundation Model Provider Risk	Foundation Model Provider Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
8	Cloud AI Provider Risk	Cloud AI Provider Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
9	Open-Source Model Risk	Open-Source Model Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.

S.NO.	TOPIC	EXPLANATION
10	Open-Source Package Risk	Open-Source Package Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
11	AI Supply Chain Risk	AI Supply Chain Risk focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
12	Contractual AI Controls	Contractual AI Controls manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
13	AI Service-Level Agreements	AI Service-Level Agreements manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
14	Data Processing Agreements	Data Processing Agreements manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.
15	Vendor Monitoring	Vendor Monitoring provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
16	Vendor Exit Strategy	Vendor Exit Strategy manages risks created when AI capabilities, models, data, infrastructure, or services depend on external parties. It includes due diligence, contractual controls, security and privacy requirements, monitoring, concentration analysis, and contingency planning.

19 AI Incident Management

S.NO.	TOPIC	EXPLANATION
1	AI Incident Detection	AI Incident Detection is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
2	AI Incident Classification	AI Incident Classification is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
3	AI Incident Reporting	AI Incident Reporting provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
4	AI Incident Response	AI Incident Response is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
5	AI Incident Escalation	AI Incident Escalation is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
6	AI Security Incident Response	AI Security Incident Response protects AI systems, models, data, prompts, APIs, and supporting infrastructure against unauthorized access, misuse, compromise, or disruption. It should use layered preventive, detective, and responsive controls based on threat modeling and least-privilege principles.
7	Model Failure Management	Model Failure Management is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
8	Hallucination Incident Management	Hallucination Incident Management concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
9	Data Breach Management	Data Breach Management is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.

S.NO.	TOPIC	EXPLANATION
10	Bias Incident Management	Bias Incident Management addresses whether AI outcomes are equitable across people, groups, and relevant attributes. It involves detecting unjustified differences, measuring them with suitable metrics, understanding their causes, and applying mitigation and monitoring controls.
11	Regulatory Notification	Regulatory Notification ensures that AI systems and related processes meet applicable laws, regulations, standards, contracts, and internal policies. It involves mapping requirements to controls, collecting evidence, monitoring changes, testing compliance, and remediating gaps.
12	Root Cause Analysis	Root Cause Analysis is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
13	Corrective Actions	Corrective Actions is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
14	Preventive Actions	Preventive Actions is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.
15	Post-Incident Review	Post-Incident Review is part of the process for detecting, containing, investigating, correcting, and learning from AI-related failures or security events. It should define severity, ownership, escalation, evidence collection, root-cause analysis, remediation, and follow-up monitoring.

20 AI Lifecycle Management

S.NO.	TOPIC	EXPLANATION
1	AI Use Case Intake	AI Use Case Intake is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
2	AI Feasibility Assessment	AI Feasibility Assessment is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
3	AI Risk Assessment	AI Risk Assessment focuses on identifying and managing the possibility that AI-related decisions, technologies, data, or dependencies could cause harm or business loss. It is managed through assessment, controls, ownership, monitoring, mitigation, and documented acceptance of residual risk.
4	Data Preparation	Data Preparation is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
5	Model Development	Model Development is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
6	Model Validation	Model Validation evaluates whether an AI system behaves as intended and remains safe, reliable, secure, and compliant under expected and challenging conditions. Testing should use defined scenarios, measurable acceptance criteria, documented evidence, and remediation for identified weaknesses.
7	Model Approval	Model Approval is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
8	Model Deployment	Model Deployment is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
9	Model Monitoring	Model Monitoring provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.

S.NO.	TOPIC	EXPLANATION
10	Model Change Management	Model Change Management is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
11	Model Retraining	Model Retraining concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.
12	Model Versioning	Model Versioning provides traceable information about AI assets, decisions, versions, dependencies, data, and changes over time. Good records improve reproducibility, accountability, impact analysis, audits, incident investigation, and regulatory evidence.
13	Model Retirement	Model Retirement is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
14	AI System Decommissioning	AI System Decommissioning is a lifecycle control governing how AI moves from idea and design through development, deployment, operation, change, and retirement. Defined criteria, approvals, documentation, testing, ownership, and monitoring help maintain control at every stage.
15	Continuous Governance	Continuous Governance defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.

21 AI Security Tools and Frameworks

S.NO.	TOPIC	EXPLANATION
1	Microsoft Presidio	Microsoft Presidio is an open-source framework for detecting and anonymizing sensitive information such as personally identifiable information. It can identify, mask, redact, or transform sensitive content before or after LLM processing.
2	Guardrails AI	Guardrails AI is a framework for validating and constraining LLM inputs and outputs. It helps enforce schemas, policies, content rules, and quality checks around generative AI applications.
3	NVIDIA NeMo Guardrails	NVIDIA NeMo Guardrails is a toolkit for defining conversational and safety controls around LLM applications. It can constrain dialog flows, reduce unsafe behavior, and add topical, security, and response guardrails.
4	Giskard	Giskard is an AI testing and evaluation platform that can identify weaknesses such as hallucination, bias, robustness issues, and security problems. It supports systematic quality and risk testing for machine-learning and generative-AI systems.
5	Promptfoo	Promptfoo is an evaluation and red-teaming framework for testing prompts, models, and LLM application behavior. It can compare outputs, run automated assertions, and test security scenarios such as prompt injection and jailbreak attempts.
6	Langfuse	Langfuse is an observability and evaluation platform for LLM applications. It helps trace prompts, model calls, latency, cost, user sessions, and quality signals so teams can monitor and improve production AI systems.
7	Rebuff	Rebuff is a framework focused on detecting and mitigating prompt-injection attacks against LLM applications. It can combine heuristics, model-based detection, and defensive checks before potentially malicious prompts reach the model.
8	OWASP LLM Guidance	OWASP LLM Guidance is an important component of an AI governance, risk, compliance, or security program. It should be defined in policy, assigned to clear owners, supported by appropriate controls, and reviewed regularly to keep AI systems trustworthy and well managed.
9	Garak	Garak is an LLM vulnerability scanner designed to probe models and applications for weaknesses such as prompt injection, data leakage, hallucination, and unsafe behavior. It is useful for automated adversarial testing and red-team exercises.
10	PyRIT	PyRIT is an open-source framework for automating security red teaming of generative-AI systems. It supports orchestrated attacks, prompt variation, scoring, and repeatable testing of model defenses.
11	ART - Adversarial Robustness Toolbox	ART - Adversarial Robustness Toolbox concerns the quality and dependability of an AI model's behavior over time and across different inputs. Teams should measure it with suitable metrics, define thresholds, investigate degradation, and retrain, recalibrate, restrict, or retire the model when necessary.

S.NO.	TOPIC	EXPLANATION
12	ModelScan	ModelScan is a security tool for scanning machine-learning model files for potentially unsafe or malicious content. It helps reduce supply-chain risk before models are loaded into trusted environments.
13	LLM Guard	LLM Guard is a security and safety toolkit for inspecting prompts and model responses. It can help detect sensitive data, malicious instructions, unsafe content, and other risks at input and output boundaries.
14	OpenAI Moderation	OpenAI Moderation refers to safety classification capabilities used to identify potentially harmful content categories in user or model-generated text. It can be incorporated as one layer within a broader content-safety and policy-enforcement architecture.
15	Cloud AI Guardrails	Cloud AI Guardrails is a guardrail used to constrain AI inputs, outputs, or actions so behavior remains within defined safety, quality, security, and compliance boundaries. It should have clear rules, fallback behavior, logging, measurable effectiveness, and periodic testing.
16	AI Observability Platforms	AI Observability Platforms provides ongoing visibility into AI performance, risk, compliance, security, and operational behavior after deployment. It uses indicators, logs, alerts, dashboards, and review cycles to detect problems and trigger timely action.
17	AI Governance Platforms	AI Governance Platforms defines how an organization directs, oversees, and controls AI-related activities. It establishes responsibilities, decision rights, policies, review mechanisms, and escalation paths so AI is used consistently and accountably.

518 topics. 21 domains. One shared reference.

From governance and model risk to generative AI security and agentic AI oversight, this guide brings a consistent vocabulary to how your organization discusses, assesses, and manages AI. Use it to brief teams, structure reviews, and anchor policy conversations in a shared language.

AI Governance, Risk, Compliance & Security — Comprehensive Topic Reference